

Symbol Tables & Binary Search Trees

Nicholas Moore

Dept. of Computing and Software, McMaster University, Canada

Acknowledgments: Material mainly based on the textbook Algorithms by Robert Sedgewick and Kevin Wayne
(Chapters 3.1, 3.2)

Symbol Table

A **symbol table** is a data structure for key-value pairs that supports two operations:

- Insert a new pair into the table (set).
- Search for the value associated with a given key (get).

Also known as: maps, dictionaries, associative arrays.

Symbol tables are **generalizes arrays** – Keys need not be between 0 and $N - 1$.

Language support: Numerous languages support symbols tables either as external libraries, built-in libraries or built-into the language (such as Python!).

Tabula Rasa

Examples:

- DNS Lookup

- key \mapsto domain name
- value \mapsto IP address

- Dictionary

- key \mapsto word
- value \mapsto definition

- Compiler

- key \mapsto variable name
- value \mapsto type

domain name	IP address
www.cs.princeton.edu	128.112.136.11
www.princeton.edu	128.112.128.15
www.yale.edu	130.132.143.21
www.harvard.edu	128.103.060.55
www.simpsons.com	209.052.165.60

↑
key value

Many, many more examples exist!

Symbol Table API

Associative array/Symbol Table abstraction. Associate one value with each key.

public class ST<Key, Value>	
ST()	<i>create a symbol table</i>
void put(Key key, Value val)	<i>put key-value pair into the table (remove key from table if value is null)</i>
Value get(Key key)	<i>value paired with key (null if key is absent)</i>
void delete(Key key)	<i>remove key (and its value) from table</i>
boolean contains(Key key)	<i>is there a value paired with key?</i>
boolean isEmpty()	<i>is the table empty?</i>
int size()	<i>number of key-value pairs in the table</i>
Iterable<Key> keys()	<i>all the keys in the table</i>

API for a generic basic symbol table

Symbol Table Conventions

Symbol table conventions adopted in the text book:

- Neither Keys nor Values are permitted to be null.
- Method `get()` returns null if key not present.
- Method `put()` overwrites old value with new value.

Intended consequences of Value \neq null

- It makes it easy to implement `contains()`.

```
public boolean contains(Key key)
{   return get(key) != null; }
```

- It allows a lazy version of `delete()`.

```
public void delete(Key key)
{   put(key, null); }
```

Ordered vs Unordered Symbol Tables

Symbol tables can be more or less generic, and as we know well, we can often improve algorithms by adding properties to data structures.

- In its most basic version, we only need a test of equality between keys.
 - Item lookup would thus become a linear search.
 - In principle this is how Python does it, since keys of different data types may be combined in one dictionary.
- If inequality operators are defined for the key data type, we can construct an **ordered symbol table**.
- We can introduce many useful operations, and improve runtimes of existing ones.
- The price we pay is key monotyping.

Ordered symbol table API

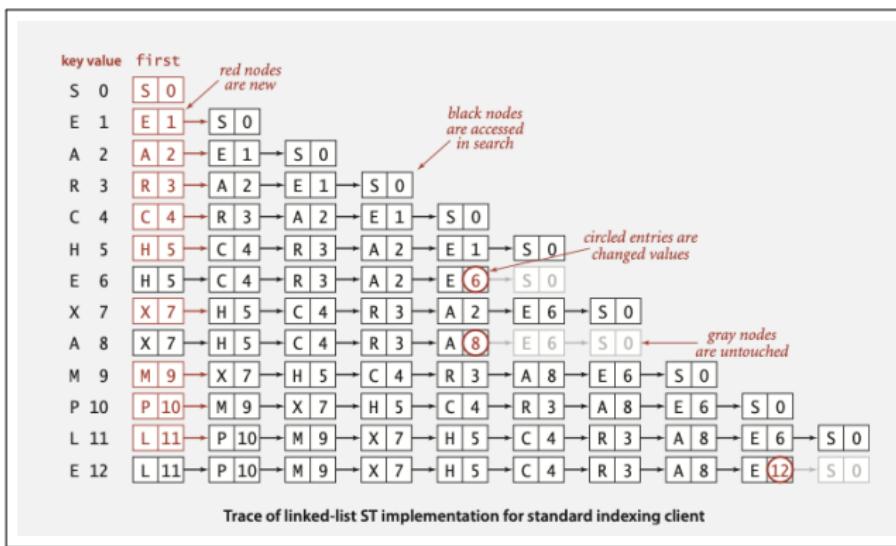
public class ST<Key extends Comparable<Key>, Value>	
ST()	<i>create an ordered symbol table</i>
void put(Key key, Value val)	<i>put key-value pair into the table (remove key from table if value is null)</i>
Value get(Key key)	<i>value paired with key (null if key is absent)</i>
void delete(Key key)	<i>remove key (and its value) from table</i>
boolean contains(Key key)	<i>is there a value paired with key?</i>
boolean isEmpty()	<i>is the table empty?</i>
int size()	<i>number of key-value pairs</i>
Key min()	<i>smallest key</i>
Key max()	<i>largest key</i>
Key floor(Key key)	<i>largest key less than or equal to key</i>
Key ceiling(Key key)	<i>smallest key greater than or equal to key</i>
int rank(Key key)	<i>number of keys less than key</i>
Key select(int k)	<i>key of rank k</i>
void deleteMin()	<i>delete smallest key</i>
void deleteMax()	<i>delete largest key</i>
int size(Key lo, Key hi)	<i>number of keys in [lo..hi]</i>
Iterable<Key> keys(Key lo, Key hi)	<i>keys in [lo..hi], in sorted order</i>
Iterable<Key> keys()	<i>all keys in the table, in sorted order</i>

API for a generic ordered symbol table

ST implementation - Unordered Linked List

Search: All nodes must be searched sequentially.

Insert: Search for the key. If present, overwrite the value, otherwise prepend a new pair to the list.

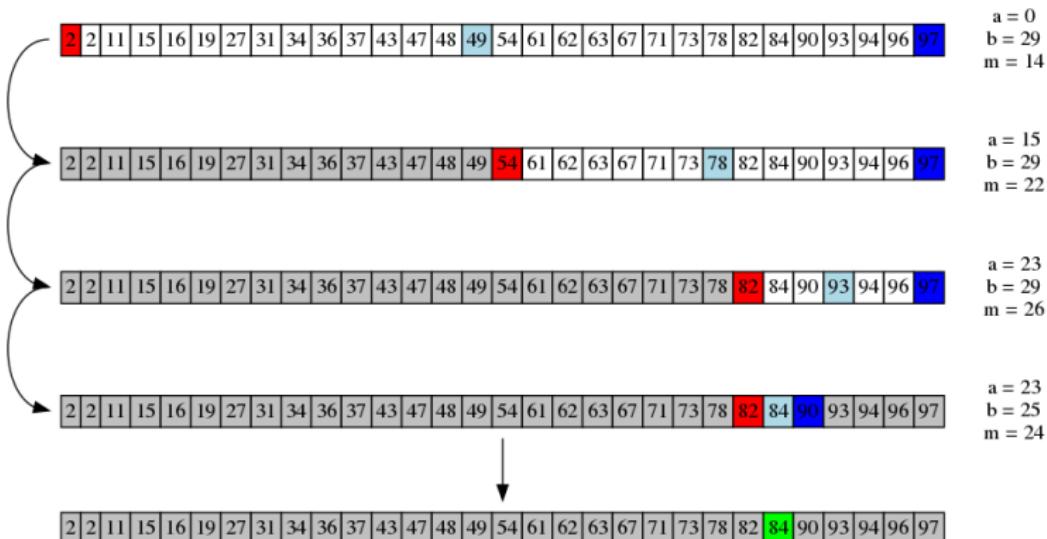


Binary Search!

Maintaining a sorted table enables the powerful **Binary Search** algorithm. Let's say we are looking for x in A :

- 1 Create two variables to keep track of the search range:
 - a to track the lower bound
 - b to track the upper bound
- 2 Examine the number at index $m = \lfloor \frac{a+b}{2} \rfloor$
 - 1 If the numbers at m , a , or b equal to x , we have found x !
 - 2 If $m < x$, we know that the index of x must be at a greater than than $\lfloor \frac{a+b}{2} \rfloor$.
 - Set a to $\lfloor \frac{a+b}{2} \rfloor + 1$ and return to step 2.
 - 3 If $m > x$, we know that x must be at a lower index than $\lfloor \frac{a+b}{2} \rfloor$.
 - Set b to $\lfloor \frac{a+b}{2} \rfloor - 1$ and return to step 2.

Visualizing Binary Search

Binary Search for $x = 84$ 

Binary Search - Java Implementation

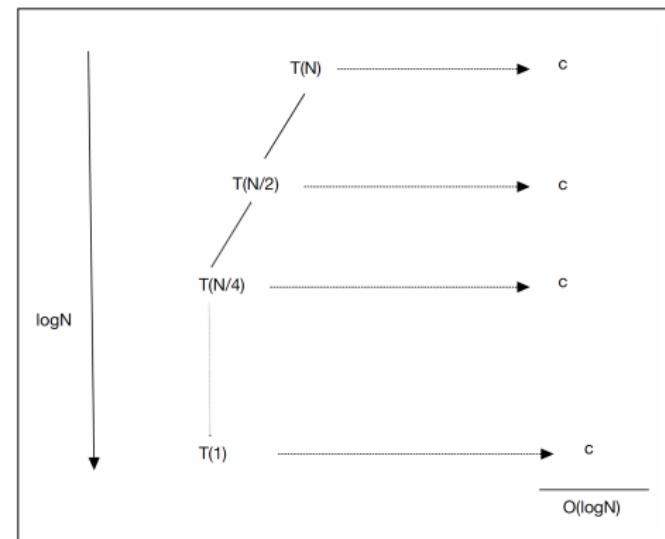
```
int a = 0;
int b = A.length - 1;
while (A <= b) { // Key is in a[a..b] or not present.
    int m = a + (b - a) / 2;
    if (x < A[m]) b = m - 1;
    else if (x > A[m]) a = m + 1;
    else return m;
}
return -1;
```

Binary Search Complexity

Binary search is a **bisection method**, as the size of the searchable section is divided by 2 at each step. Our recurrence equation is therefore

$$T(N) = T(N/2) + c, \text{ with } c > 0 \text{ and } T(1) = 1$$

- For simplicity assume N is a power of 2.
- From the recursion tree (\Rightarrow), we have $T(N) = 1 + c \log_2 N$
- Therefore $T(N) \in O(\log N)$
- Binary search uses at most $1 + \log_2 N$ key comparisons to search a sorted array.



Ordered ST: Insert

Unfortunately, insertion into the middle of an array requires that all items greater than the item inserted be shifted one place to the right.

keys[]										vals[]												
key	value	0	1	2	3	4	5	6	7	8	9	N	0	1	2	3	4	5	6	7	8	9
S	0	S										1	0									
E	1	E	S									2	1	0								
A	2	A	E	S								3	2	1	0							
R	3	A	E	R	S							4	2	1	3	0						
C	4	A	C	E	R	S						5	2	4	1	3	0					
H	5	A	C	E	H	R	S					6	2	4	1	5	3	0				
E	6	A	C	E	H	R	S					6	2	4	6	3	0					
X	7	A	C	E	H	R	S	X				7	2	4	6	5	3	0	7			
A	8	A	C	E	H	R	S	X				7	8	4	6	5	3	0	7			
M	9	A	C	E	H	M	R	S	X			8	8	4	6	5	9	3	0	7		
P	10	A	C	E	H	M	P	R	S	X		9	8	4	6	5	9	10	3	0	7	
L	11	A	C	E	H	L	M	P	R	S	X	10	8	4	6	5	11	9	10	3	0	7
E	12	A	C	E	H	L	M	P	R	S	X	10	8	4	12	5	11	9	10	3	0	7
		A	C	E	H	L	M	P	R	S	X		8	4	12	5	11	9	10	3	0	7

Symbol table (ordered and unordered) operations summary

Sequential Search (unordered linked list):

- Search: $O(N)$
- Insert: $O(N)$

Binary Search (ordered array):

- Binary Search: $O(\lg N)$
- Insert: $O(N)$

In Summary...

underlying data structure	implementation	pros	cons
<i>linked list (sequential search)</i>	SequentialSearchST	best for tiny STs	slow for large STs
<i>ordered array (binary search)</i>	BinarySearchST	optimal search and space, order-based ops	slow insert

Binary Search Trees

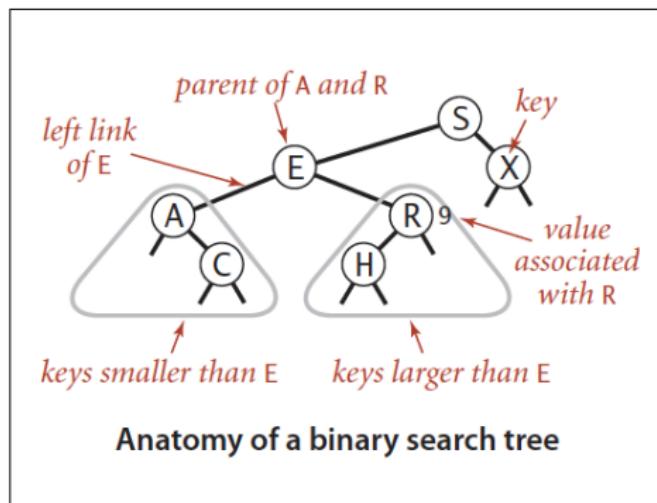
A **Binary Search Tree (BST)** is a binary tree where each node has a key.

Each key is:

- larger than all keys in its left subtree
- smaller than all keys in its right subtree

In contrast to binary heaps, *BSTs need not be complete binary trees*.

- Later on we'll have to address the problems this introduces.



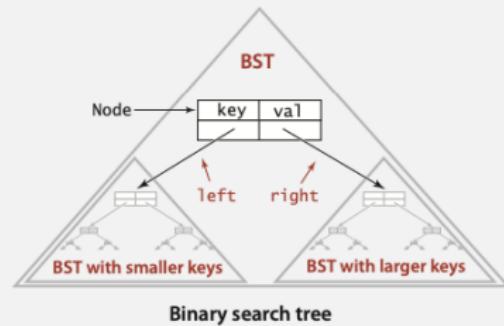
If we're using this to implement a symbol table, can you have duplicate keys in the BST? What about in general?

BST implementation: Node

A BST **Node** is composed of four fields:

- A Key and a Value.
- A reference to the left (smaller) and right (larger) subtree.
- (For later) An instance variable N that gives the node count in the subtree rooted at the node. This field facilitates the implementation of various ordered symbol-table operations

```
private class Node
{
    private Key key;
    private Value val;
    private Node left, right;
    public Node(Key key, Value val)
    {
        this.key = key;
        this.val = val;
    }
}
```



Key and Value are generic types; Key is Comparable

BST Skeleton

```
public class BST<Key extends Comparable<Key>, Value>
{
    private Node root;                                ← root of BST

    private class Node
    { /* see previous slide */ }

    public void put(Key key, Value val)
    { /* see next slides */ }

    public Value get(Key key)
    { /* see next slides */ }

    public void delete(Key key)
    { /* see next slides */ }

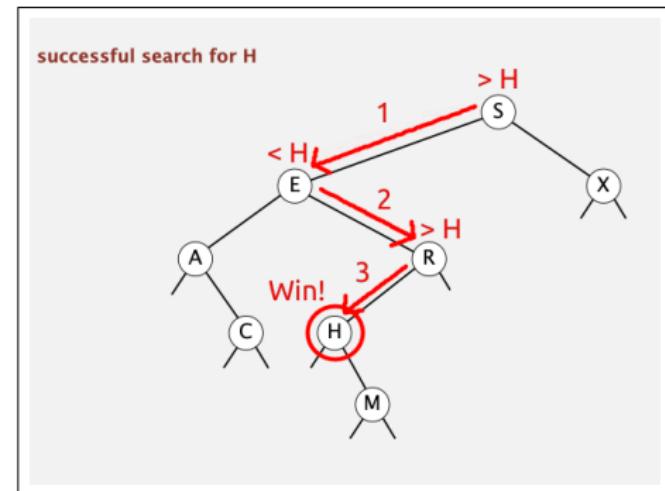
    public Iterable<Key> iterator()
    { /* see next slides */ }

}
```

Binary Search Tree – Find/Search

BST Search Procedure (k is the searched for key):

- If the current node's key is greater than k , recurse on the left branch.
- If the current node's key is less than k , recurse on the right branch.
- If the current node's key is equal to k , return the value.
- If the branch you're recursing on is empty, the search fails!



BST Implementation: get()

Get: Return value corresponding to given key, or null if no such key. We look for the *key* starting from the *root* node, and do the below for each node.

Get algorithm outline:

- If $key = node.key$ return node's value.
- If $key < node.key$ recurse on the left subtree.
- If $key > node.key$ recurse on the right subtree.

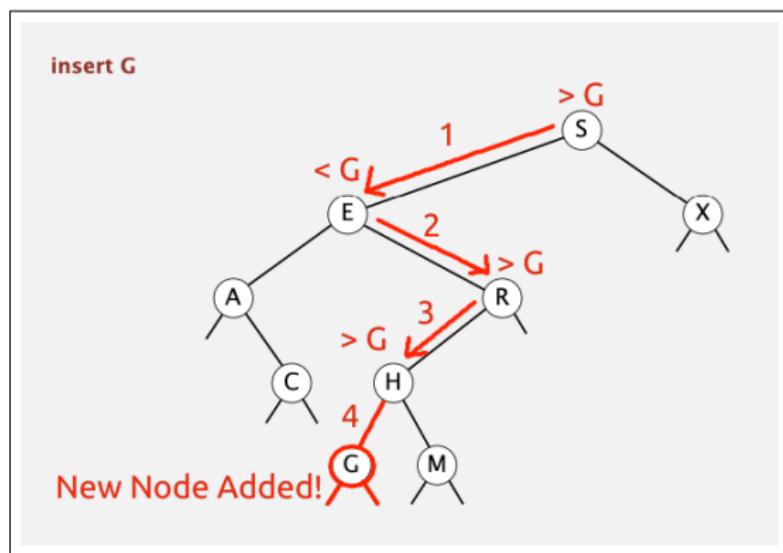
```
public Value get(Key key)
{
    Node x = root;
    while (x != null)
    {
        int cmp = key.compareTo(x.key);
        if      (cmp < 0) x = x.left;
        else if (cmp > 0) x = x.right;
        else if (cmp == 0) return x.val;
    }
    return null;
}
```

Cost: Number of comparisons is equal to $1 + \text{depth of node}$.

Binary Search Tree – Insert

Same as search procedure, except for what happens when you reach a null branch.

- Rather than an empty branch indicating failure, this is where we insert the new node.
- The structure of a BST is *dependent on the order items are added!*
- The best case scenario is a complete binary tree, which is unlikely to happen naturally.



BST insert() / put()

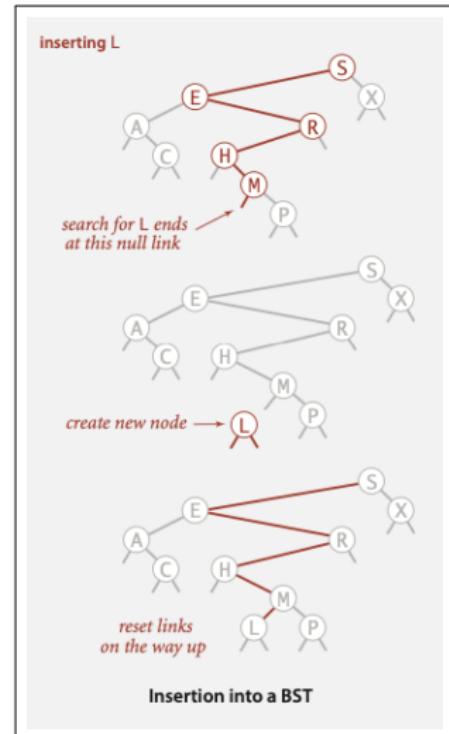
Put: Associates a key with a value.

Search for the key.

- If the key is in tree, overwrite the value.
- If the key is not in the tree, add a new node for it.

Implementation:

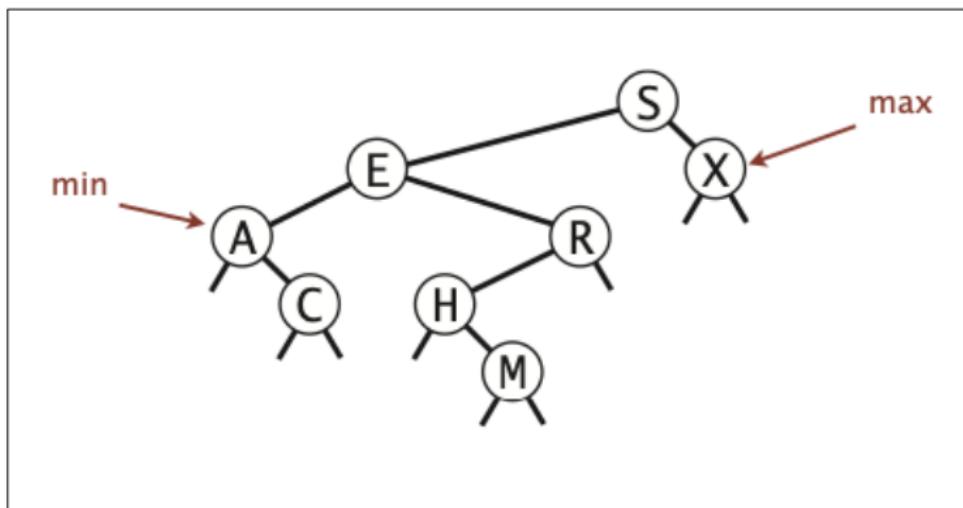
- Can be recursive or iterative (similar to get())
- **Cost:** Number of comparisons is equal to $1 + \text{the depth of node}$.



BST: Min. and Max. Operations

Minimum - returns the smallest key in table. Go to the left as far as possible.

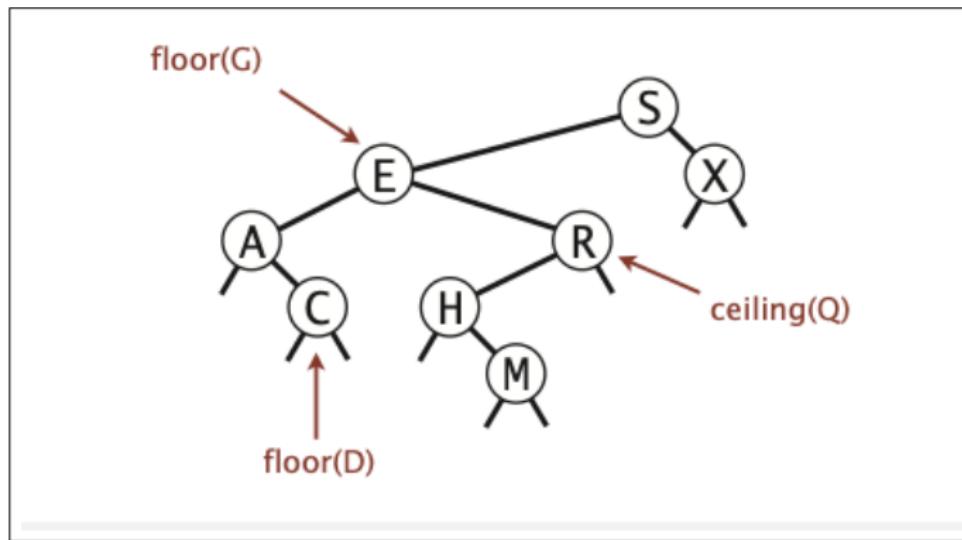
Maximum - returns the largest key in table. Go to the right as far as possible.



BST: Floor and Ceiling Operations

Floor - the largest key in the BST less than or equal to the key we are flooring.

Ceiling - the smallest key in the BST greater than or equal to the key we are cieling...ing...



Case 1. [k equals the key in the node]

The floor of k is k .

Case 2. [k is less than the key in the node]

The floor of k is in the left subtree.

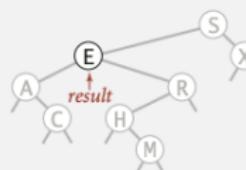
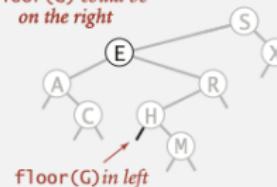
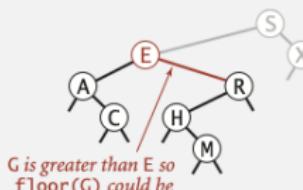
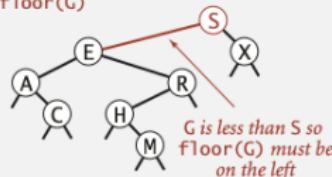
Case 3. [k is greater than the key in the node]

The floor of k is in the right subtree

(if there is any key $\leq k$ in right subtree);

otherwise it is the key in the node.

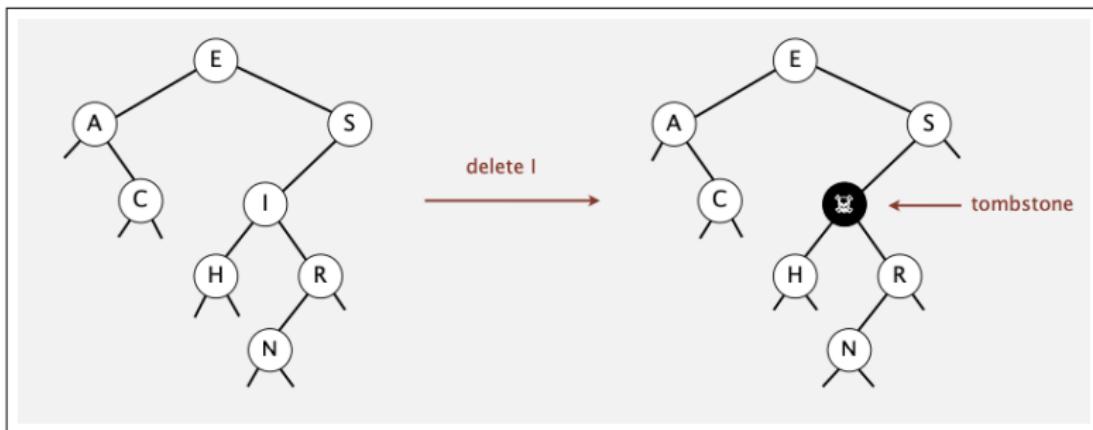
finding $\text{floor}(G)$



Lazy Deletion

To remove a node with a given key:

- Set its value to null.
- Leave key in tree to guide search (but don't consider it equal in search).

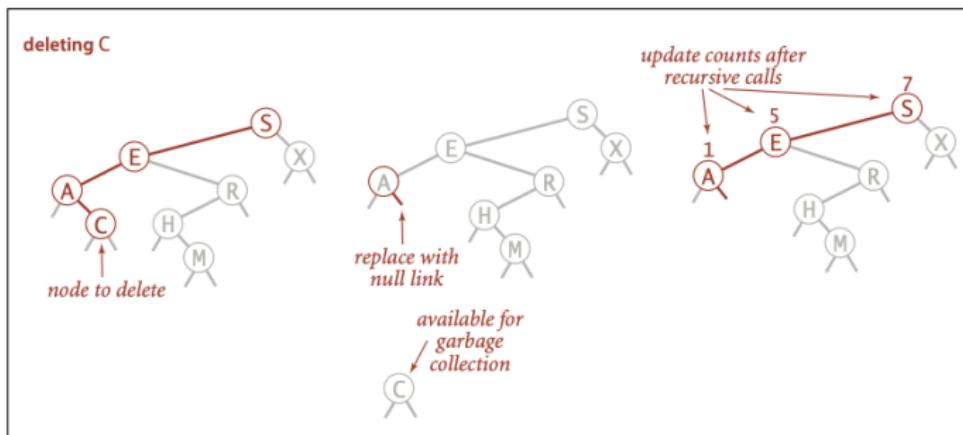


Unsatisfactory solution. Tombstones occupy memory!

BST: Hibbard Deletion

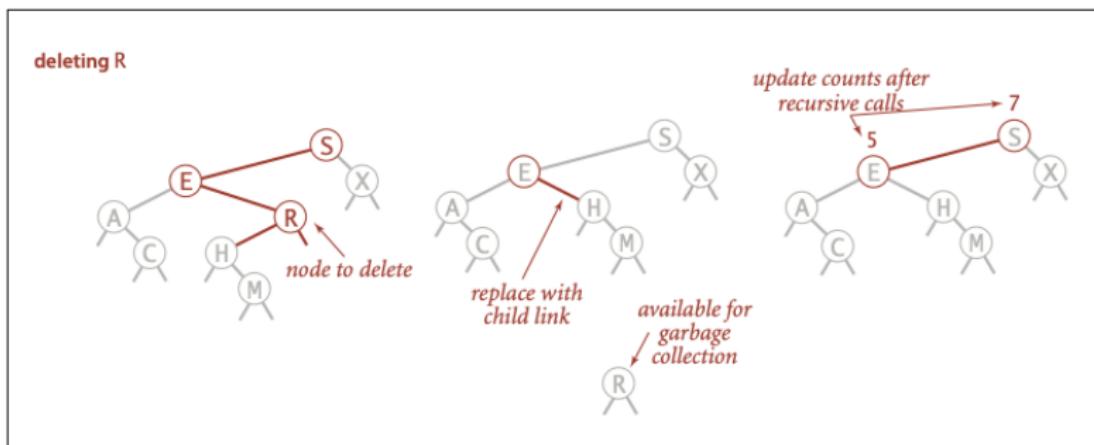
To delete a node with key k : search for node t containing key k .

Case 0: [0 children] Delete t by setting its parent link to null.



BST: Hibbard Deletion

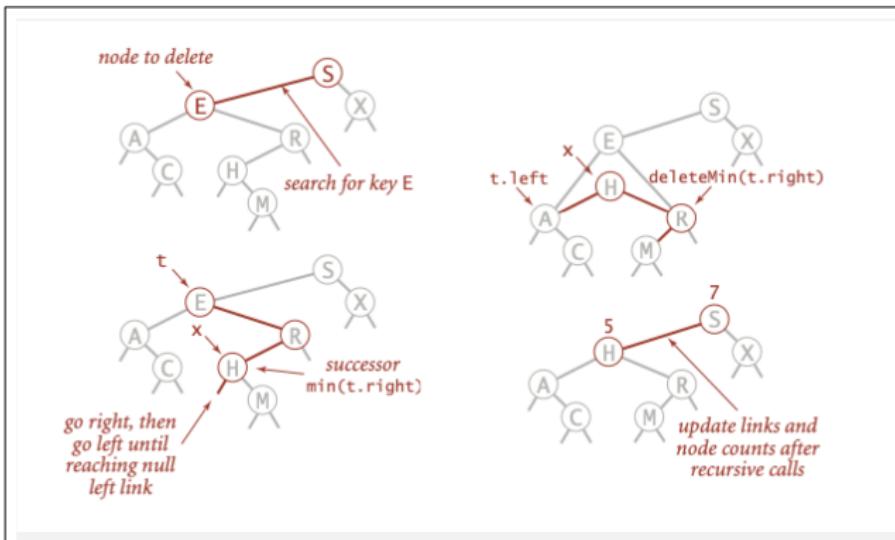
Case 1: [1 child] Delete t and connect its single child to t 's parent.



BST: Hibbard Deletion:

Case 2. [2 children]

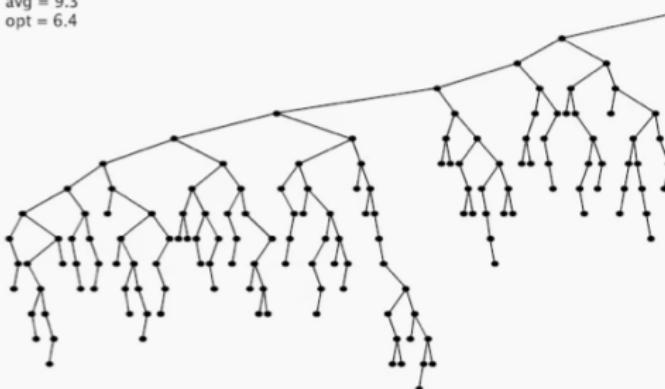
- t is replaced by x = the minimum key in t's right subtree.
- x's right child is x's replacement.
- x. left = t.left, x.right=t.right (if x = t.right, then x.right =null).



BST: Hibbard deletion analysis

Unsatisfactory solution. Not symmetric.

N = 150
max = 16
avg = 9.3
opt = 6.4



Surprising consequence. Trees not random (!) $\Rightarrow \sqrt{N}$ per op.

Longstanding open problem. Simple and efficient delete for BSTs.

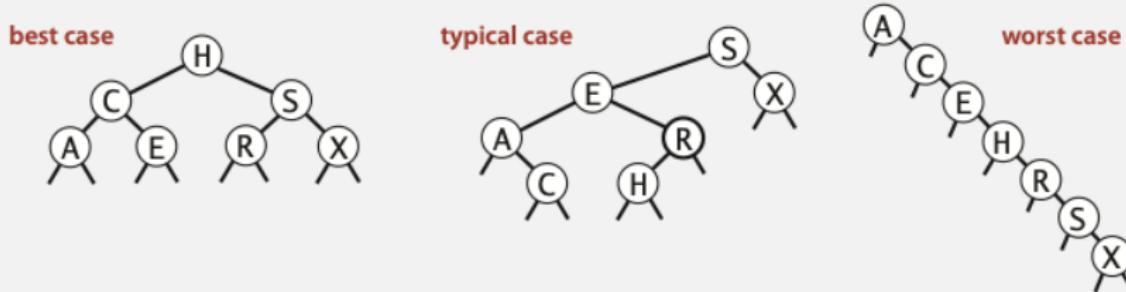
BST Cost

algorithm (data structure)	worst-case cost (after N inserts)		average-case cost (after N random inserts)		efficiently support ordered operations?
	search	insert	search hit	insert	
<i>sequential search (unordered linked list)</i>	N	N	N/2	N	no
<i>binary search (ordered array)</i>	$\lg N$	N	$\lg N$	N/2	yes
<i>binary tree search (BST)</i>	N	N	$1.39 \lg N$	$1.39 \lg N$	yes

Cost summary for basic symbol-table implementations (updated)

BST Tree Shape

- One set of keys can be stored in many differently structured BSTs.
- Remember: the number of comparisons for search/insert is proportional to the depth of the tree!



Tree shape, *and therefore runtime*, depends on the order of insertion! Not fantastic!

BST Tree Randomization

Assume that the keys inserted in a uniform random order.

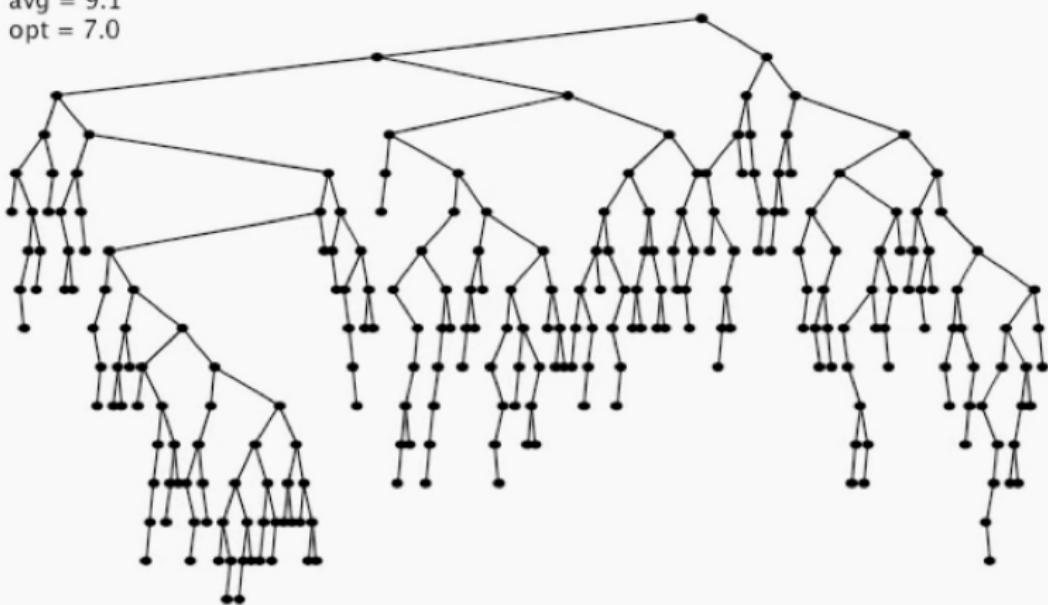
- That is, the probability that any remaining node is the next node to be added is a **uniform probability distribution**.

In a BST built from N random keys:

- Search hits/misses and insertions about $1.39 \log_2 N$ comparisons on average.

BST insertion: random order visualization

$N = 255$
 $\max = 16$
 $\text{avg} = 9.1$
 $\text{opt} = 7.0$



Tree Balancing Algorithms

There are several tree balancing algorithms, some of which are beyond the scope of this course:

- T-Tree - Used in main-memory databases such as MySQL
- Treap - Randomizes tree structure with every insertion.
- **Red-Black Tree** - Nodes are dynamically “colored” red or black, which informs insert procedures.
- **B-Tree** - Generalizes BSTs to allow nodes with more than 2 children. Good for file systems.
- **2-3 Tree** - Specific type of B-Tree, where nodes either have 2 children and one datum or 3 children and two data.